

THE EVOLUTION OF LMOS

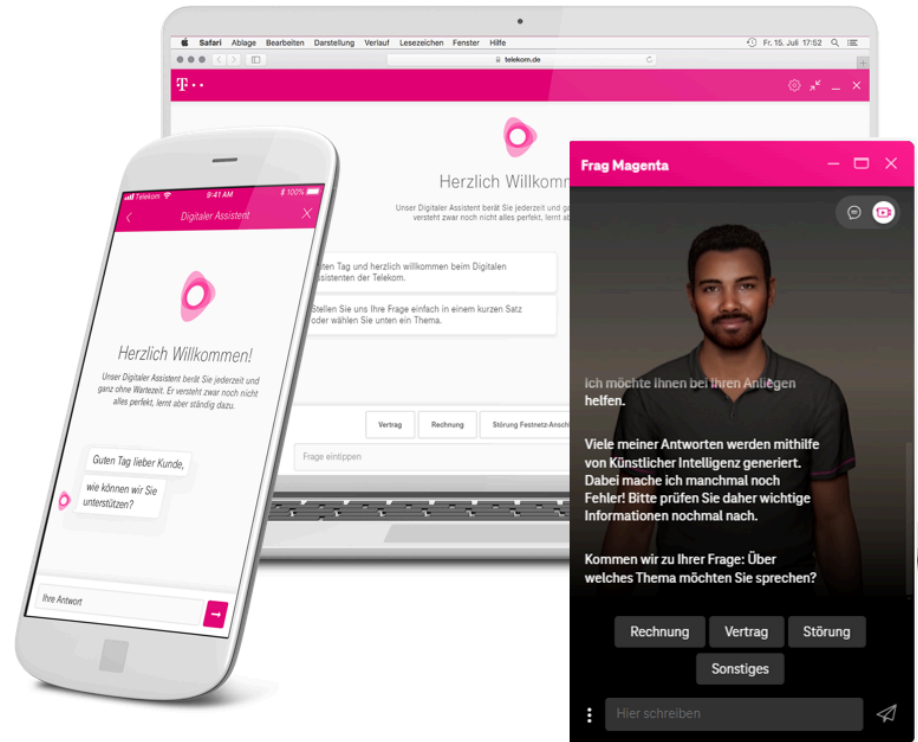
ROBERT WINKLER



LIFE IS FOR SHARING.

Frag Magenta

- Digital assistant of DT Customer Service
- Available for Web, OneApp, WhatsApp, Apple Business Chat and IVR in Germany



Facing limitations

RASA's Design:

- **Scripted Dialogue Flow:** Predefined YAML-based scripts to steer dialogues.
- **Predictive Planning Required:** Customer interactions must be anticipated in advance.
- **Intent Classification Challenges:** Natural Language Understanding (NLU) struggles with accurately classifying user intents.
- **Manual Effort:** Continuous effort is required to update intents and train NLU.
- **Knowledge gap:** No FAQ knowledge base was available.
- **Scripting:** YAML is no scripting/programming language.

```
stories:  
- story: beginning of flow  
  steps:  
  - intent: greet  
  - action: action_ask_user_question  
  - checkpoint: check_asked_question  
  
- story: handle user affirm  
  steps:  
  - checkpoint: check_asked_question  
  - intent: affirm  
  - action: action_handle_affirmation  
  - checkpoint: check_flow_finished  
  
- story: handle user deny  
  steps:  
  - checkpoint: check_asked_question  
  - intent: deny  
  - action: action_handle_denial  
  - checkpoint: check_flow_finished  
  
- story: finish flow  
  steps:  
  - checkpoint: check_flow_finished  
  - intent: goodbye  
  - action: utter_goodbye
```

Overcoming limitations

Key Pain Points:

- **High Maintenance Costs:** Updating scripted dialogues is time-consuming and labor-intensive.
- **Customer Frustration:** Unresolved inquiries due to lack of knowledge led to negative customer experiences.
- **Call Center Costs:** Increased volume of customer queries at the call center, driving up costs.

Need for Change:

- **Improve NLU/NLP:** The existing NLU/NLP, static scripts and knowledge sources were not sufficient for many customer queries.
- **Objectives:** Increase solution rate, reduce call center volumes, lower maintenance costs, and enhance customer satisfaction through improved NLU/NLP, increased knowledge and less manual scripting efforts.

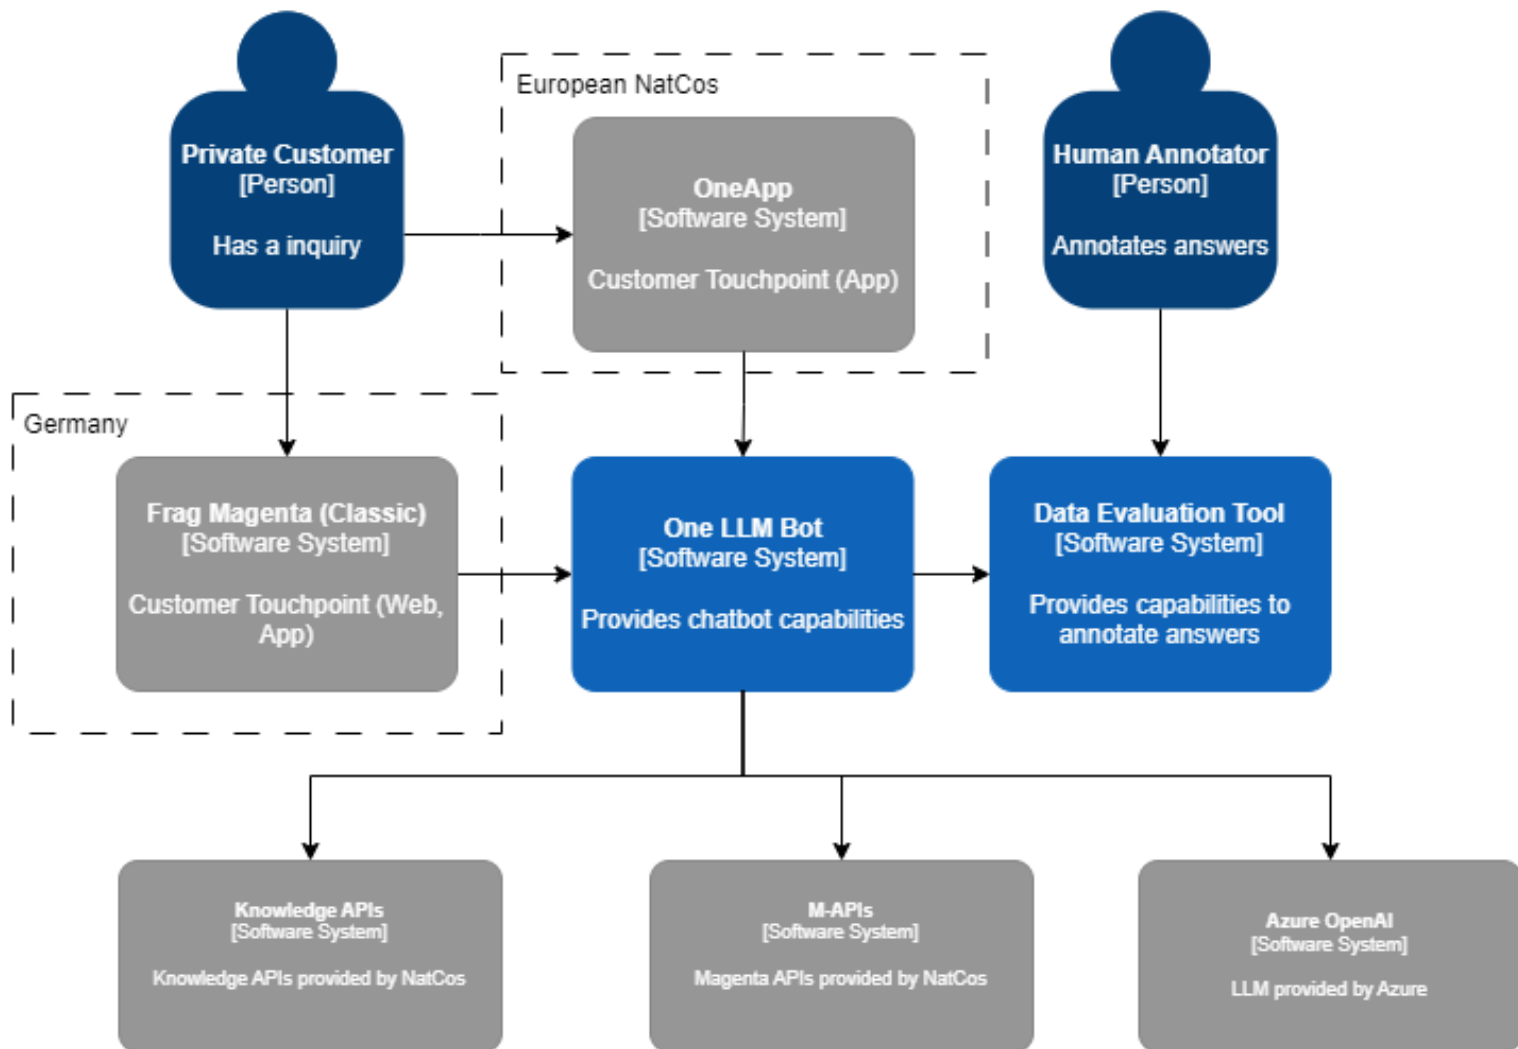
Taking a bold decision

Advancements in Generative AI:

- **Large Language Models (LLMs):** Models showed potential for better language understanding and processing.
- **Opportunity:** Use LLMs and Retrieval Augmented Generation (RAG) to improve NLU/NLP and reasoning/planning.

Decision to Innovate:

- **Multi-tenant and omni-channel system:** A strategic move to develop a single system for multiple NatCos and channels by making use of M-APIs.
- **Multi-agent system (MAS):** Multiple LLM-based agents, each focused on a specific business domain, working together to solve customer inquiries.
- **Highly configurable:** Every tenant can have a unique set of Agents, capabilities and knowledge sources.
- **Objective:** Improve speed to rollout the digital assistant to multiple NatCos. Starting with Germany, Austria and Croatia.



Facing new territory and doubts

Industry Landscape:

- **Focus on Single-Agent Systems:** Most companies/frameworks were focused on single-agent RAG solutions.
- **Lack of Frameworks:** No established, production-ready multi-agent frameworks were available.

Team Concerns:

- **Distributed Team:** Team is distributed about multiple countries: Germany, India, Greece, ...
- **Expertise Gap:** Dev team's background in Java/Kotlin rather than AI-specific technologies/languages.
- **Complexity:** Concerns about handling the complexity of a multi-tenant, multi-agent system and reaching the efficiency needed for Deutsche Telekom.

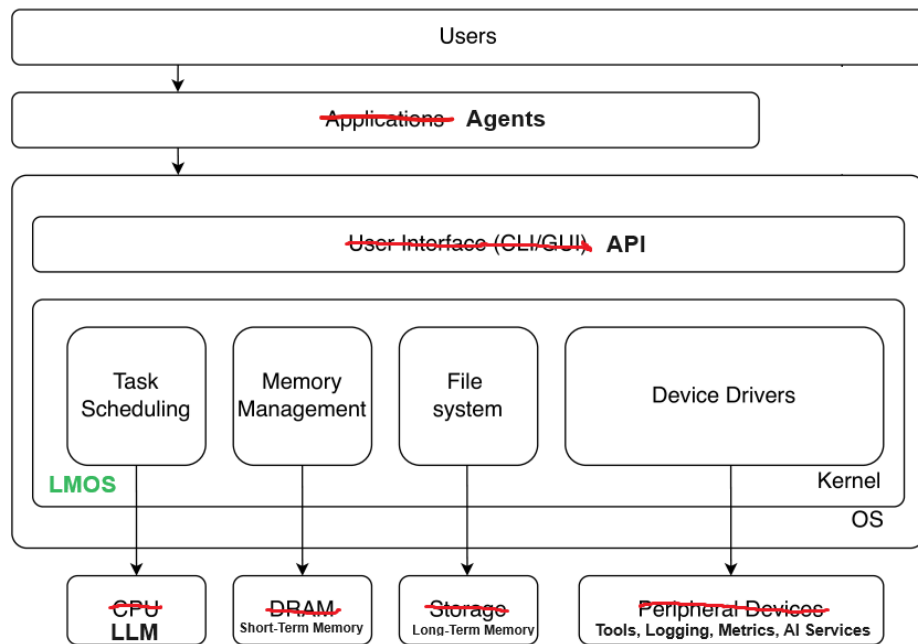
Internal Doubts in LT:

- **Building vs. Adapting:** Debate over whether to create a new system or adapt/buy existing solutions.
- **Feasibility Concerns:** Doubts about the team's ability to deliver a new scalable and efficient solution in time.

Finding Inspiration (1/3)

Operating System Design:

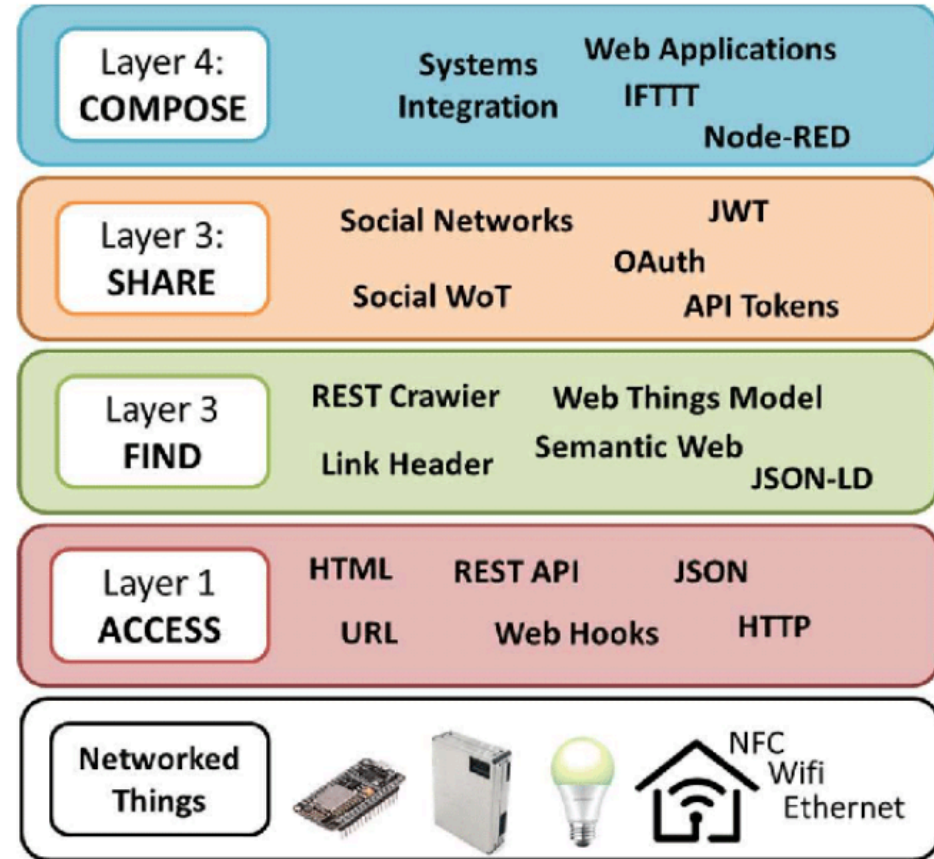
- **OS/Kernel Analogy:** LMOS acts as a Kernel/OS between Agents and Infrastructure providing an abstraction layer.
- **Application Development:** An OS provides system calls --> LMOS provides an API to simplify agent development.
- **Abstractions:** An OS abstracts hardware --> LMOS abstracts the complexities of working with LLMs, memory and tools.
- **Application Management:** An OS manages applications --> LMOS manages Agents and is doing execution planning.



Finding Inspiration (2/3)

Web of Things -> Web of Agents:


- **Standardized Communication (Access):**
Built on top of open web standards and data models.
- **Discoverability and extendability (Find):**
Dynamic discoverability of self-describing agents, allowing the multi-agent system to evolve.
- **Reusability (Share):** A single agent can be reused by multiple tenants concurrently.
- **Cross-Platform Compatibility (Compose):**
Freedom to compose agents without being locked into a single platform.



Finding Inspiration (3/3)

SpaceX Engineering Model:

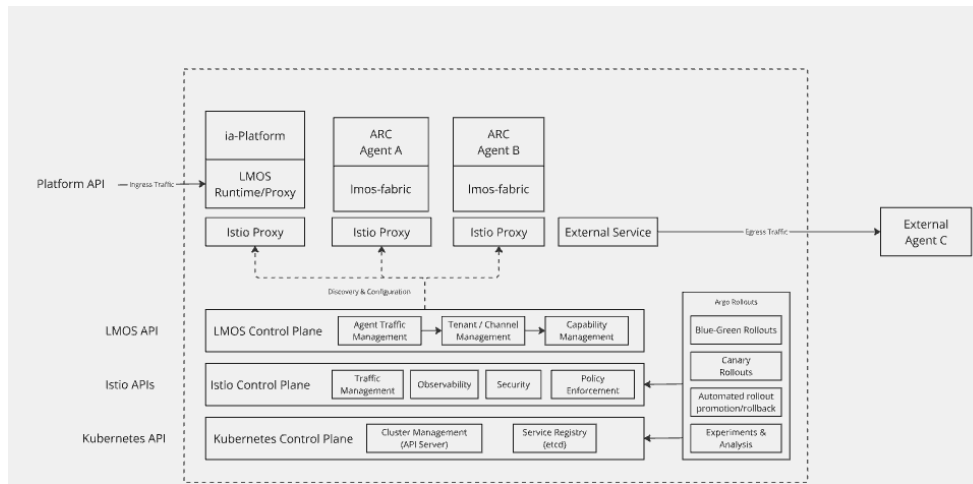
- **Frequent Deployments:** Frequent Agent releases with real-world feedback for quick improvement.
- **Parallelism:** Reusable launch system --> Develop and test agents simultaneously for faster innovation.
- **Reuseability:** Reusable rocket system --> LMOS provides reusable modules, e.g. Flows and Steps.
- **Reuseability:** Reusable rockets --> Reuse agents across multiple tenants.



The foundation

Technical Foundation:

- **Built on Kubernetes:** Leverage Kubernetes for orchestration and scalability.
- **Istio Service Mesh:** Utilize Istio as the service mesh to enhance traffic management, security, and observability across agents.
- **Extend Kubernetes:** Extend Kubernetes capabilities by developing a custom control plane to manage agents.
- **ArgoCD & GitOps:** Implemented ArgoCD with a GitOps approach to automate deployments and perform canary releases.



The concepts

Key Concepts:

- **Agent Registry:** Agents register their meta-data in a central registry.
- **Agent Runtime:** A runtime responsible for orchestrating the collaboration between multiple AI agents.
- **Agent Discovery:** Runtime can discover installed Agents and their capabilities.
- **Dynamic Routing:** Dynamically route queries to the most suitable agent. LMOS uses advanced methods like language models and vector embeddings to intelligently match queries to the right agent.
- **Knowledge Sharing:** The Runtime ensures that agents share context, memory, and knowledge as needed to handle customer queries holistically.
- **Memory Management:** LMOS includes built-in memory management for agents, allowing them to store and retrieve data during interactions.
- **Operator:** Listens for new or modified Channels and Agents and dynamically resolves capabilities.

Overcoming hurdles (1/2)

Development challenges:

- **Technical challenges:** Encountering various technical obstacles such as model performance and hallucinations.
- **Channel-Specific Requirements:** Addressing unique technical needs for different channels, such as voice, web or app.
- **High Testing Effort:** Demanding and resource-intensive testing processes, requiring extensive manual annotations and validations to guarantee model reliability.
- **Multilingual Data Anonymization:** Developing robust Named Entity Recognition (NER) models capable of handling data anonymization across diverse languages and linguistic structures.
- **Multilingual Language Support:** Ensuring language and embedding models possess the capability to process, comprehend, and maintain accuracy across multiple languages.

Overcoming hurdles (2/2)

International Collaboration:

- **Team Efforts:** Collaboration between teams in Germany and India.
- **Coordinated Work:** Overcame time time zone and way of working differences to ensure effective teamwork.

Pressure to Prove Concept:

- **Stakeholder Engagement:** Continuously communicate and showcase the platform's value and potential to gain the trust and support of stakeholders.

The critical test

Deployment Requirements:

- **Scalability:** Ability to handle large-scale deployment across multiple use cases and NatCos.
- **Security and Efficiency:** Ensuring platform security and operational efficiency.

Real-World Testing:

- **Performance Evaluation:** Assessment of the platform's performance and accuracy in live settings.
- **Competitiveness:** Comparison with leading industry products and solutions, such as Rasa CALM or Sprinklr.

The outcome

Live Implementation:

- **Handling Interactions:** Efficiently handled thousands of customer interactions accross multiple channels and diverse use cases.

Performance Metrics:

- **Solution Rate:** Achieved an 85% solution rate.
- **Hallucinations:** Less than 5% incorrect or irrelevant responses, demonstrating strong model reliability.
- **Development Efficiency:** Developed 14 use cases within a month, averaging 2.5 days per use case, highlighting rapid iteration and implementation.

Impact:

- **Higher solution rate:** Enhanced solution rate by leveraging additional context information, such as knowledge sources and M-APIs.
- **Accelerated Processes:** Accelerated development and deployment timelines, leading to quicker delivery of new use cases.

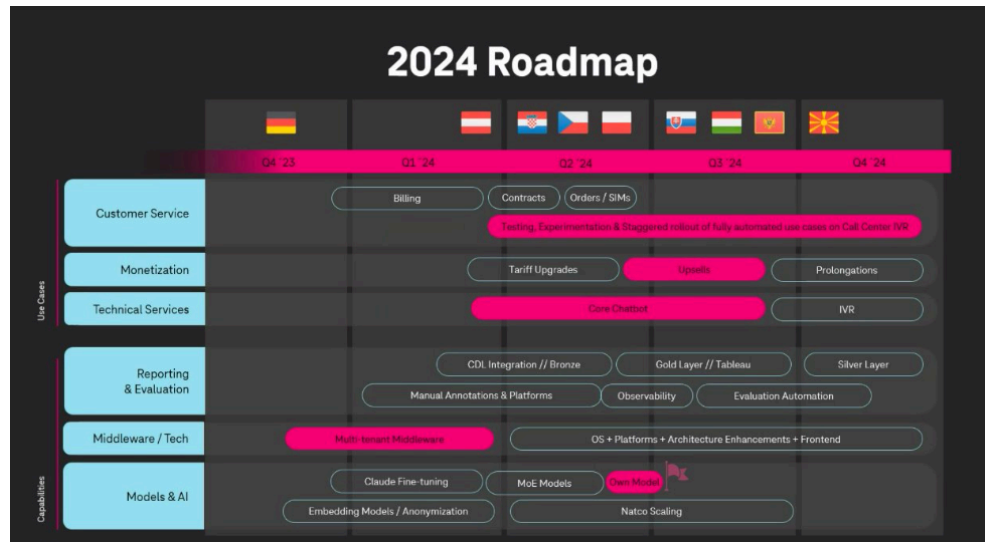
Moving forward

Business Value:

- **Return on Investment:** Delivered measurable business value - reduced call volumes and improved customer experience.

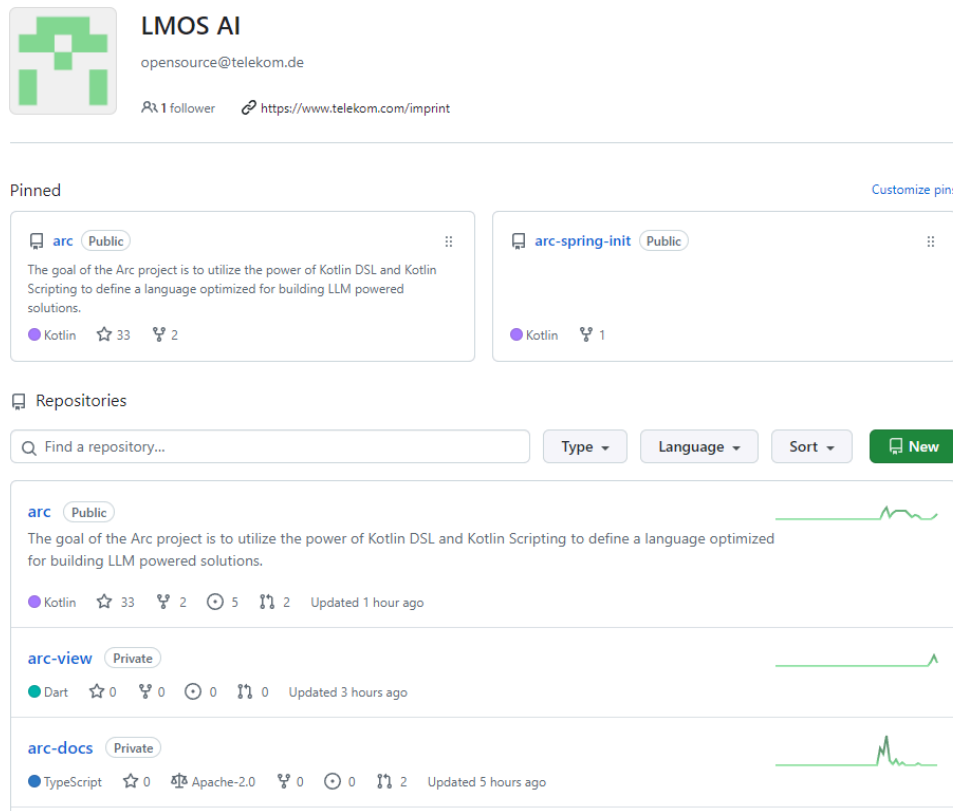
Impact on Deutsche Telekom:

- **Strategic Foundation:** LMOS became the cornerstone for Deutsche Telekom's chatbot strategy in Europe.
- **NatCo Expansion:** Successfully rolled out in Austria, with Croatia as the next country on the roadmap.
- **International Teamwork:** A success of an international team pushing beyond traditional



Going Open Source

- **Community Contribution:** Release of LMOS to the open-source community on GitHub.
- **Sharing Innovations:** Providing a valuable solution for building an enterprise-ready, multi-tenant, multi-channel and multi-agent system.



The screenshot shows the GitHub profile of LMOS AI. The profile header includes a green and white pixelated avatar, the name "LMOS AI", the email "opensource@telekom.de", and a link to the imprint page. Below the header, there are two pinned repositories: "arc" and "arc-spring-init". The "arc" repository is public, has 33 stars, 2 forks, and 5 watchers, and is written in Kotlin. The "arc-spring-init" repository is also public, has 1 fork, and is written in Kotlin. Below the pinned repositories, there is a section for all repositories, which includes "arc-view" (private, 0 stars, 0 forks, 0 watchers, written in Dart) and "arc-docs" (private, 0 stars, 0 forks, 2 watchers, written in TypeScript). The "arc" repository description states: "The goal of the Arc project is to utilize the power of Kotlin DSL and Kotlin Scripting to define a language optimized for building LLM powered solutions."

LMOS AI
opensource@telekom.de
1 follower <https://www.telekom.com/imprint>

Pinned

arc Public
The goal of the Arc project is to utilize the power of Kotlin DSL and Kotlin Scripting to define a language optimized for building LLM powered solutions.
Kotlin 33 2 5

arc-spring-init Public
Kotlin 1

Repositories

Find a repository... Type Language Sort New

arc Public
The goal of the Arc project is to utilize the power of Kotlin DSL and Kotlin Scripting to define a language optimized for building LLM powered solutions.
Kotlin 33 2 5 Updated 1 hour ago

arc-view Private
Dart 0 0 0 Updated 3 hours ago

arc-docs Private
TypeScript 0 Apache-2.0 0 0 2 Updated 5 hours ago

Questions?